

DATASCI W261: Machine Learning at Scale

HW6.6

Angela Gunn & Jing Xu
angela@egunn.com
jaling@gmail.com

MIDS - W261 - Section 3
February 24, 2016
Assignment : Week 6

Introduction

This is a map-reduce version of expectation maximization algo for a mixture of Gaussians model. There are two mrJob MR packages, mr_GMixEmIterate and mr_GMixEmInitialize. The driver calls the mrJob packages and manages the iteration.

The multivariate Gaussian probability density function is defined as follows:

Equation 1:

$$\mathcal{N}(x|\mu, \Sigma) = \frac{1}{(2\pi)^{\frac{1}{2}} |\Sigma|^{\frac{1}{2}}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$

Expectation Step:

Given priors, mean vector and covariance matrix, calculate the probability of that each data point belongs to a class

$$\omega_k^i = p(w_k | x^i, \theta) \quad \left| \begin{array}{l} \text{equation 1} \\ \Rightarrow \end{array} \right. = \frac{\pi_k \mathcal{N}(x^i | \mu_k, \Sigma_k)}{\sum_{j=1}^K \mathcal{N}(x^i | \mu_j, \Sigma_j)}$$

where

- k = number of classes
- i = example index
- θ = (phi, mu, sigma)
- μ = centroid
- Σ = covariance matrix
- \mathcal{N} = number of samples
- π = prior class probability

$\mathcal{N}(x^i | \mu_k, \Sigma_k)$ is calculated as below

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right) \quad \left| \begin{array}{l} \text{equation 2} \\ \Rightarrow \end{array} \right.$$

Where $|\Sigma|$ is the determinant of covariance matrix

M Step:

$$\text{Centroid}_k \Rightarrow \hat{\mu}_k = \frac{1}{n_k} \sum_{i=1}^n p(w_k | x_i, \theta) x^i$$

which is calculated as

$$\text{Centroid}_k \Rightarrow \mu_k \quad \left| \begin{array}{l} \text{equation 3} \\ \Rightarrow \end{array} \right. = \frac{\sum_{i=1}^n \omega_k^i x^i}{\sum_{i=1}^m \omega_k^i}$$

$$\left. \begin{aligned} \text{Covariance}_k &\Rightarrow \sum_k^{\wedge} \\ &= \frac{1}{n_k} \sum_{i=1}^n p(w_k|x_i, \theta) (x^i \\ &- \hat{\mu}_k)(x^i - \hat{\mu}_k)^T \end{aligned} \right|$$

which is calculated as

$$\left. \begin{aligned} \text{Covariance}_k &\Rightarrow \sum_k \\ &= \frac{\sum_{i=1}^n \omega_k^i (x^i - \mu_k)(x^i - \mu)^T}{\sum_{i=1}^n \omega_k^i} \end{aligned} \right| \text{equation 4}$$

$$\text{Prior}_k \Rightarrow \hat{\pi}_k = \frac{n_k}{n} \text{ where } n_k = \sum_{i=1}^n p(w_k|x_i, \theta)$$

which is calculated as

$$\left. \begin{aligned} \text{Prior} &\Rightarrow \phi_k = \frac{1}{n} \sum_{i=1}^n \omega_k^i \end{aligned} \right| \text{equation 5}$$

\Rightarrow

Data Generation

```
In [3]: %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
import json
size1 = size2 = size3 = 1000 #set size 1000

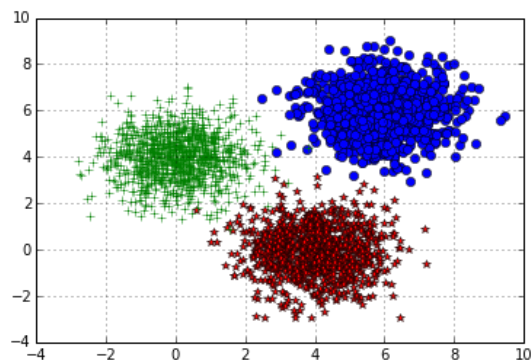
#generate random arrays for our clusters; append to data
samples1 = np.random.multivariate_normal([4, 0], [[1, 0],[0, 1]], size1)
data = samples1
samples2 = np.random.multivariate_normal([6, 6], [[1, 0],[0, 1]], size2)
data = np.append(data,samples2, axis=0)
samples3 = np.random.multivariate_normal([0, 4], [[1, 0],[0, 1]], size3)
data = np.append(data,samples3, axis=0)

# Randomize data
data = data[np.random.permutation(size1+size2+size3),]
with open("data.txt", "w") as f:
    for row in data.tolist():
        json.dump(row, f)
        f.write("\n")
```

Data Visualization

```
In [5]: #plot
plt.plot(samples1[:, 0], samples1[:, 1], '*', color = 'red' )
plt.plot(samples2[:, 0], samples2[:, 1], 'o', color = 'blue')
plt.plot(samples3[:, 0], samples3[:, 1], '+', color = 'green')

plt.grid(True)
plt.show()
```



Initialization

Initialization

Here suppose we know there are 3 components

```
In [7]: %%writefile mr_GMM_Init.py
from mrjob.job import MRJob
from mrjob.step import MRStep

from numpy import mat, zeros, shape, random, array, zeros_like, dot, linalg
from random import sample
import json
from math import pi, sqrt, exp, pow

class GMM_Init(MRJob):
    DEFAULT_PROTOCOL = 'json'

    def steps(self):
        return [
            MRStep(mapper=self.mapper,
                    reducer=self.reducer)
        ]

    def __init__(self, *args, **kwargs):
        #call base init method
        super(GMM_Init, self).__init__(*args, **kwargs)

        self.numMappers = 1      #number of mappers
        self.count = 0

    def configure_options(self):
        #call base configure_options method
        super(GMM_Init, self).configure_options()

        self.add_passthrough_option(
            '--k', dest='k', default=3, type='int',
            help='k: number of densities in mixture')
        self.add_passthrough_option(
            '--pathName', dest='pathName', default="", type='str',
            help='pathName: pathname where intermediateResults.txt is stored')

    def mapper(self, key, values):
        #something simple to grab random starting point
        #collect the first 2k points (this isn't 2000, this is 2*the number of densities)
        #The values are points
        if self.count <= 2*self.options.k:
            self.count += 1
            yield (1,values)

    def reducer(self, key, values):
        #receives an integer as the key, and a data point as the values.

        #accumulate data points mapped to 0 from 1st mapper and pull out k of them as starting point
        centroids = []

        #load the values and append to centroids; then emit 1, and the value
        for xj in values:
            x = json.loads(xj)
            centroids.append(x)
            yield 1, xj

        #sample the points - these are our new centroids
        index = sample(range(len(centroids)), self.options.k)
        centroids_sample = []
        for i in index:
            centroids_sample.append(centroids[i])

        #use the covariance of the selected centers as the starting guess for covariances

        #first, calculate mean of centers
        #create an array of the centroids, then divide each by the number of clusters(k)
        mean_array = array(centroids_sample[0])
        for i in range(1,self.options.k):
            mean_array = mean_array + array(centroids_sample[i])

        mean_array = mean_array/float(self.options.k)
```

```

        #then accumulate the deviations
        covariance = zeros((len(mean_array),len(mean_array)),dtype=float)
        for x in centroids_sample:
            deviation = array(x) - mean_array
            for i in range(len(mean_array)):
                covariance[i,i] = covariance[i,i] + deviation[i]*deviation[i]

        covariance = covariance/(float(self.options.k))
        covariance_inverse = linalg.inv(covariance)

        covariance_inverse1 = [covariance_inverse.tolist()]*self.options.k

        #for debugging
        jDebug = json.dumps([centroids_sample,mean_array.tolist(),covariance.tolist(),covariance_inverse.t
olist(),covariance_inverse1])
        debugPath = self.options.pathName + 'debug_init.txt'
        fileOut = open(debugPath,'w')
        fileOut.write(jDebug)
        fileOut.close()

        #also need a starting guess at the phi's - prior probabilities
        #initialize them all with the same number - 1/k - equally probably for each cluster

        phi = zeros(self.options.k,dtype=float)

        for i in range(self.options.k):
            phi[i] = 1.0/float(self.options.k)

        #form output object
        outputList = [phi.tolist(), centroids_sample, covariance_inverse1]

        jsonOut = json.dumps(outputList)

        # Write to file
        fullPath = self.options.pathName + 'intermediateResults.txt'
        with open(fullPath, 'w') as outfile:
            outfile.write(jsonOut)

if __name__ == '__main__':
    GMM_Init.run()

```

Overwriting mr_GMM_Init.py

mrJob Class File

```

In [8]: %%writefile mr_GMM_Iterate.py
from mrjob.job import MRJob

from math import sqrt, exp, pow,pi
from numpy import zeros, shape, random, array, zeros_like, dot, linalg
import json

def gauss(x, mu, P_1):
    #function for calculating gaussian
    xtemp = x - mu
    n = len(x)
    p = exp(- 0.5*dot(xtemp,dot(P_1,xtemp)))
    detP = 1/linalg.det(P_1)
    p = p/(pow(2.0*pi,n/2.0)*sqrt(detP))
    return p

class GMM_Iterate(MRJob):
    DEFAULT_PROTOCOL = 'json'

    def __init__(self, *args, **kwargs):
        #call original init function
        super(GMM_Iterate, self).__init__(*args, **kwargs)

        fullPath = self.options.pathName + 'intermediateResults.txt'

        #read input
        fileIn = open(fullPath)
        inputJson = fileIn.read()
        """

```

```

fileIn.close()
inputList = json.loads(inputJson)

#break out the individual lists:

#prior class probabilities
temp = inputList[0]
self.phi_array = array(temp)
#current means list
temp = inputList[1]
self.mean_array = array(temp)
#inverse covariance matrices for w, calc.
temp = inputList[2]
self.covariance_inverse1_array = array(temp)

#accumulate partial sums
#sum of weights - by cluster
self.new_phi_array = zeros_like(self.phi_array)      #partial weighted sum of weights
self.new_means_array = zeros_like(self.mean_array)
self.new_cov_array = zeros_like(self.covariance_inverse1_array)

self.numMappers = 1          #number of mappers
self.count = 0               #passes through mapper

def configure_options(self):
    #call base configure_options
    super(GMM_Iterate, self).configure_options()

    self.add_passthrough_option(
        '--k', dest='k', default=3, type='int',
        help='k: number of densities in mixture')
    self.add_passthrough_option(
        '--pathName', dest='pathName', default="", type='str',
        help='pathName: pathname where intermediateResults.txt is stored')

def mapper(self, key, values):
    #accumulate partial sums for each mapper
    xList = json.loads(values)
    x = array(xList)

    #E Step
    #Estimate class assignments(responsibilities) or weights.
    #Use the current model to estimate class assignments.

    #weighted vectors
    wtVect = zeros_like(self.phi_array)
    for i in range(self.options.k):
        wtVect[i] = self.phi_array[i]*gauss(x,self.mean_array[i],self.covariance_inverse1_array[i])
    wtSum = sum(wtVect)
    wtVect = wtVect/wtSum

    #accumulate to update est of probability densities.
    #increment count
    self.count += 1

    #M STEP
    #For each cluster K, update the centroid running
    #We are calculating the summation parts of the formulas

    #accumulate weights for phi est (the Prior)
    self.new_phi_array = self.new_phi_array + wtVect

    for i in range(self.options.k):
        #accumulate weighted x's for mean calc (the Centroid)
        self.new_means_array[i] = self.new_means_array[i] + wtVect[i]*x

        #accumulate weighted squares for cov estimate (the Covariance)
        xmm = x - self.mean_array[i]
        covInc_array = zeros_like(self.new_cov_array[i])

        for l in range(len(xmm)):
            for m in range(len(xmm)):
                covInc_array[l][m] = xmm[l]*xmm[m]
        self.new_cov_array[i] = self.new_cov_array[i] + wtVect[i]*covInc_array

```

```

        self.new_cov_array[i] = self.new_cov_array[i] + newcov[i].covariance_array
        #dummy yield - real output passes to mapper_final in self

def mapper_final(self):
    #output the new arrays
    out = [self.count, (self.new_phi_array).tolist(), (self.new_means_array).tolist(), (self.new_cov_array).tolist()]
    jOut = json.dumps(out)

    yield 1,jOut

def reducer(self, key, xs):
    #accumulate partial sums
    first = True

    #accumulate partial sums
    #xs is a list of partial stats, including count, phi, mean, and covariance.

    #Each stats is k-length array, storing info for k components
    for val in xs:
        if first:
            temp = json.loads(val)
            #totCount, totPhi, totMeans, and totCov are all arrays
            totCount = temp[0]
            totPhi_array = array(temp[1])
            totMeans_array = array(temp[2])
            totCov_array = array(temp[3])
            first = False
        else:
            temp = json.loads(val)
            #cumulative sum of four arrays
            totCount = totCount + temp[0]
            totPhi_array = totPhi_array + array(temp[1])
            totMeans_array = totMeans_array + array(temp[2])
            totCov_array = totCov_array + array(temp[3])

    #At this point, we take our Centroid, Covariance and Prior and divide by sums collected from mapper

    #Prior is divided by totCount
    #Centroid & Covariance divided by the n_k (prior)

    #finish calculation of new probability parameters. array divided by array
    #This is the PRIORS
    newPhi_array = totPhi_array/totCount

    #initialize these to get the right size arrays
    #Means = CENTROIDS
    #Cov = COVARIANCE
    newMeans_array = totMeans_array
    newCov_1_array = totCov_array
    for i in range(self.options.k):
        newMeans_array[i,:] = totMeans_array[i,:]/totPhi_array[i]
        tempCov_array = totCov_array[i,:,:]/totPhi_array[i]
        #almost done. just need to invert the cov matrix. invert here to save doing a matrix inversion

        #with every input data point.
        newCov_1_array[i,:,:] = linalg.inv(tempCov_array)

    outputList = [newPhi_array.tolist(), newMeans_array.tolist(), newCov_1_array.tolist()]
    jsonOut = json.dumps(outputList)

    #write new parameters to file
    fullPath = self.options.pathName + 'intermediateResults.txt'
    fileOut = open(fullPath,'w')
    fileOut.write(jsonOut)
    fileOut.close()

if __name__ == '__main__':
    GMM_Iterate.run()

```

Overwriting mr_GMM_Iterate.py

Driver and Execution

```
In [11]: from mr_GMM_Init import GMM_Init
```

```

In [11]: from mr_GMM_Init import GMM_Init
from mr_GMM_Iterate import GMM_Iterate
import json
from math import sqrt
import matplotlib.pyplot as plt

def plot_iteration(means):

    plt.plot(samples1[:, 0], samples1[:, 1], '.', color = 'c')
    plt.plot(samples2[:, 0], samples2[:, 1], '.', color = 'c')
    plt.plot(samples3[:, 0], samples3[:, 1], '.', color = 'c')

    for b,g,r in means:
        plt.plot(b[0], b[1], 'o', markersize = 10, color = 'b')

        plt.plot(g[0], g[1], 'o', markersize = 10, color = 'g')

        plt.plot(r[0], r[1], 'o', markersize = 10, color = 'r')

    b,g,r = means[-1]
    plt.plot(b[0], b[1], '*', markersize = 10, color = 'w')
    plt.plot(g[0], g[1], '*', markersize = 10, color = 'w')
    plt.plot(r[0], r[1], '*', markersize = 10, color = 'w')

    b,g,r = means[0]
    plt.plot(b[0], b[1], '*', markersize = 10, color = 'black')
    plt.plot(g[0], g[1], '*', markersize = 10, color = 'black')
    plt.plot(r[0], r[1], '*', markersize = 10, color = 'black')

    plt.title("Movement of Centroids (black = start, white = final)")

    plt.show()

def dist(x,y):
    #euclidean distance between two lists
    sum = 0.0
    for i in range(len(x)):
        temp = x[i] - y[i]
        sum += temp * temp
    return sqrt(sum)

#first run the initializer to get starting centroids
filePath = 'data.txt'
myPathName = '/home/hduser/Notebooks/'
mrJobInit = GMM_Init(args=[filePath, '--pathName', myPathName, '--no-strict-protocol'])

with mrJobInit.make_runner() as runner:
    runner.run()

#pull out the centroid values to compare with values after one iteration
emPath = "intermediateResults.txt"
fileIn = open(emPath)
paramJson = fileIn.read()
fileIn.close()

delta = 10
iter_num = 0
centroids = []
#Begin iteration on change in centroids
while delta > 0.02:
    print "Iteration " + str(iter_num)
    iter_num = iter_num + 1
    #parse old centroid values
    oldParam = json.loads(paramJson)
    #run one iteration
    oldMeans = oldParam[1]
    mrJob2 = GMM_Iterate(args=[filePath, '--pathName', myPathName, '--no-strict-protocol'])
    with mrJob2.make_runner() as runner:
        runner.run()

    #compare new centroids to old ones
    fileIn = open(emPath)
    paramJson = fileIn.read()
    fileIn.close()
    newParam = json.loads(paramJson)

    k_means = len(newParam[1])
    newMeans = newParam[1]

```

```

delta = 0.0
for i in range(k_means):
    delta += dist(newMeans[i],oldMeans[i])
#print oldMeans
strOut = "delta={:>5.3f}".format(delta)
for x,y in oldMeans:
    strOut = "{}\t[{:>7.4f},{:>7.4f}]".format(strOut, x, y)
print strOut
centroids.append(oldMeans)
#plot_iteration(oldMeans)

#plot_iteration(oldMeans)
print "Iteration" + str(iter_num)
strOut = "delta={:>5.3f}".format(delta)
for x,y in newMeans:
    strOut = "{}\t[{:>8.4f},{:>8.4f}]".format(strOut, x, y)
print strOut
centroids.append(newMeans)
plot_iteration(centroids)
print "\n***FINISHED***"

```

```

Iteration 0
delta=4.721    [ 1.5503, 3.2216]    [-1.3424, 5.3565]    [ 4.8742, 5.4802]
Iteration 1
delta=0.771    [ 2.7086, 1.2657]    [ 0.1186, 4.7070]    [ 5.5891, 5.9377]
Iteration 2
delta=0.456    [ 2.5747, 1.3355]    [ 0.0394, 4.3397]    [ 5.8326, 5.9516]
Iteration 3
delta=0.358    [ 2.7322, 1.1616]    [-0.0494, 4.2986]    [ 5.9547, 5.9711]
Iteration 4
delta=0.319    [ 2.9487, 0.9509]    [-0.0744, 4.2666]    [ 5.9697, 5.9738]
Iteration 5
delta=0.307    [ 3.1515, 0.7582]    [-0.0927, 4.2334]    [ 5.9707, 5.9745]
Iteration 6
delta=0.315    [ 3.3501, 0.5707]    [-0.0998, 4.2005]    [ 5.9704, 5.9747]
Iteration 7
delta=0.317    [ 3.5496, 0.3803]    [-0.0904, 4.1629]    [ 5.9701, 5.9747]
Iteration 8
delta=0.268    [ 3.7364, 0.1991]    [-0.0604, 4.1155]    [ 5.9696, 5.9746]
Iteration 9
delta=0.161    [ 3.8797, 0.0561]    [-0.0186, 4.0664]    [ 5.9689, 5.9744]
Iteration 10
delta=0.070    [ 3.9567,-0.0239]    [ 0.0160, 4.0307]    [ 5.9685, 5.9738]
Iteration 11
delta=0.026    [ 3.9859,-0.0545]    [ 0.0348, 4.0113]    [ 5.9683, 5.9732]
Iteration 12
delta=0.009    [ 3.9956,-0.0648]    [ 0.0427, 4.0032]    [ 5.9682, 5.9728]
Iteration13
delta=0.009    [ 3.9987, -0.0682]    [ 0.0455, 4.0002]    [ 5.9682, 5.9726]

```

